

The Learning Parity with Noise (LPN) Problem, The BKW Algorithm for LPN

1 The LPN Problem

Definition 1. The LPN_n^τ problem asks to find the secret $s \in \mathbb{Z}_2^n$, given samples of the form $(a_i, b_i = \langle a_i, s \rangle \oplus e_i) \in \mathbb{Z}_2^n \times \mathbb{Z}_2$, where $a_i \xleftarrow{\$} \mathbb{Z}_2^n$, e_i follows the Bernoulli distribution with parameter τ (i.e. $\Pr(e_i = 1) = \tau$), for $0 < \tau < \frac{1}{2}$. Such e_i s are sometimes referred to as “noise”, τ being the noise rate. We define the LPN bias $j = \frac{1}{2} - \tau$.

Remark. 1. The hardness of LPN is defined by two parameters: n and τ (we know τ). In particular:

- $\tau = 0 \implies$ Gaussian elimination
- $\tau = \frac{1}{2} \implies$ e_i s act as a one-time pad. (s is information theoretically hidden)
- if $\tau > \frac{1}{2}$, $\forall i$ take $b_i \oplus 1$, you will get LPN samples with $\tau' = \tau - \frac{1}{2}$.

2. (s, e) is uniquely defined with high probability given $\mathcal{O}(n)$ LPN samples.

3. Applications of LPN:

- identification schemes ([HB01])
- collision-resistant hash functions
- encryption ([Ale03])
- light weight crypto

2 Useful bounds

Mitzenmacher & Upfal, “Probability and Computing” [MU05]

Chernoff bound. Let e_1, \dots, e_m be independant Bernoulli random variables s.t. $\forall i, \Pr(e_i = 1) = \tau$, let $N = \sum_{i=1}^m e_i$, $\mu = \mathbb{E}(N) = m\tau$. Then for $0 < \delta < 1$,

$$\Pr(N \geq \mu(1 + \delta)) \leq e^{-\frac{\mu\delta^2}{3}}$$

$$\Pr(N \leq \mu(1 - \delta)) \leq e^{-\frac{\mu\delta^2}{3}}$$

For $x \in \{0, 1\}^n$, denote the Hamming weight of x as $wt(x) = \#\{i | x_i = 1\}$.

Fact 2. For $x \in \{0, 1\}^n$ whose entries are iid from Ber_τ , $wt(x)$ follows the Binomial distribution $Bin_{n,\tau}$.

Corollary 3. Let e be the error vector coming from m -many LPN_n^τ samples. Then it holds that, for any $\delta \in]0, 1[$, $\Pr(wt(e) > m\tau(1 + \delta)) \leq e^{-\frac{m\tau\delta^2}{3}}$.

In particular, for all $\epsilon > 0$, set $\delta = m^{\frac{-\epsilon}{2}}$, $\Pr(wt(e) > \tau(m + m^{\frac{-\epsilon}{2}})) \leq e^{-\frac{\tau m^{1-\epsilon}}{3}}$

Corollary 4 (Majority vote). Fix a bit $x \in \{0, 1\}$. Let $b_i = x \oplus e_i$ where $e_i \leftarrow Ber_\tau$. Denote the bias by $\eta = \frac{1}{2} - \tau$. Then we can decide x with probability at least $1 - e^{-c}$ (where c is a constant), having $m \geq \frac{3}{2} \frac{c(1-2\eta)}{\eta^2}$ many b_i s.

Proof. $x' = \lfloor \frac{1}{m} \sum_{i=1}^m b_i \rfloor = \lfloor \frac{1}{m} \sum_{i=1}^m x + e_i \rfloor = x + \lfloor \frac{1}{m} \sum_{i=1}^m e_i \rfloor$.

If $\frac{1}{m} \sum_{i=1}^m e_i < \frac{1}{2}$ then $x' = x$.

Use Chernoff bound with $\delta = \frac{2\eta}{1-2\eta}$, then $\Pr(\frac{1}{m} wt(e) > \frac{1}{2}) < e^{-c}$ if $m \geq \frac{3}{2} \frac{c(1-2\eta)}{\eta^2}$ \square

3 Algorithms for LPN_n^τ

Assume:

- we have as many samples as we need
- the secret is unique

3.1 Brute-Force

3.1.1 Over s

Algorithm 1 Input: m -many LPN_n^τ samples

Output: s

```

1: for all  $s \in \{0, 1\}^n$  do
2:   if  $Test(s)$  then
3:     return  $s$ 
4:   end if
5: end for

```

with $Test$:

Test

Input: m -many LPN_n^τ samples $(a_i, b_i)_{i \leq n}$

s : candidate for the solution

$\epsilon > 0$

Output: Reject / Accept

```

1:  $N = 0$ 
2: for  $i \in [m]$  do
3:    $N += b_i \oplus \langle a_i, s \rangle$ 
4: end for
5: if  $N > \tau(m + \frac{1}{m^{\frac{\epsilon}{2}}})$  then
6:   Reject
7: else
8:   Accept
9: end if
```

Claim 5. Test accepts the correct s with probability $> 1 - e^{-\frac{\tau m^{1-\epsilon}}{3}}$, Test rejects a wrong s with probability $> 1 - e^{-\frac{\tau m}{3}(1-\tau)^2}$

Proof. The first part of the claim follows from corollary 1. With $s \neq s^*$, $s[1] = s^*[1]$,

$b_i \oplus \langle a_i, s \rangle = \langle a_i, s \oplus s^* \rangle + e_i = a_i[1] + e_i$.

$\Pr([a_i[1] + e_i] = 1) = \Pr(a_i[1] = 1) \cdot \Pr(e_i = 0) + \Pr(a_i[1] = 0) \cdot \Pr(e_i = 1) = \frac{1}{2}(1 - \tau) + \frac{1}{2}\tau = \frac{1}{2}$.

In this case, $N \sim \text{Bin}_{m, \frac{1}{2}}$. Apply Chernoff bound with $\tau = \frac{1}{2}$ and $\delta = 1 - 2\tau(1 + \frac{1}{m^{\frac{\epsilon}{2}+1}})$ \square

3.1.2 Over e

Algorithm 2

```

1: for all integers  $t \leq \tau(n + \frac{1}{n^{\frac{\epsilon}{2}}})$  do
2:   for all  $e' \in \{0, 1\}^n$  s.t.  $\text{wt}(e') = t$  do
3:     Solve  $A.x = b - e'$  for  $x$  // require  $A \in \mathcal{GL}_n(\mathbb{F}_2)$ 
4:     if  $\text{Test}(x)$  then
5:       return  $x$  (or  $e'$ )
6:     end if
7:   end for
8: end for
```

Theorem 6. Algorithm 1 solves the LPN_n^τ problem with high probability in time $T(\text{brute force}) = \tilde{\mathcal{O}}(2^n)$ using $m = \mathcal{O}(n)$ samples and $\mathcal{O}(n)$ memory.

Algorithm 2 solves the LPN_n^τ problem in time $T = \tilde{\mathcal{O}}(2^{\binom{n}{\tau n}}) = \tilde{\mathcal{O}}(2^{n \cdot H(\tau)})$ using $\mathcal{O}(n)$ memory samples.

3.2 “Many samples algorithm”

Algorithm 3 Goal: determine s_1

Repeat sampling until you see (a_i, b_i) where $a_i = (1, 0, \dots, 0) = \vec{e}_i // b_i = \langle a_i, s \rangle + e_i = s_1 + e_i$

Repeat until we found m -many such samples

\Rightarrow we can deduce s_1 if $m \geq \Theta(\frac{1}{(\frac{1}{2} - \tau)^2})$ //cf. Corollary 4 (Majority Vote)

This algorithm requires $m \frac{1}{\Pr(a_i = e_i)}$ many repetitions, i.e. $T = m2^n = \#\text{samples}$.

Try Gaussian elimination to obtain \vec{e}_i .

Problem: using naive Gaussian elimination, the vector \vec{e}_i appears as a sum of $\mathcal{O}(n^2)$ a_i s.

Lemma 7 (Piling-up lemma). *Let e_i be iid Bernoulli random variables s.t. $\Pr(e_i = 1) = \tau \forall i \in [m]$.*

Then $\sum_{i=1}^m e_i \sim Ber_{\frac{1}{2} - \frac{1}{2}(1-2\tau)m}$

Proof by induction. 1. $m = 1 \rightarrow \text{OK}$

2. Assume it holds for $m - 1$

3.

$$\begin{aligned} \Pr\left(\sum_{i=1}^m e_i = 1\right) &= \Pr\left(\sum_{i=1}^{m-1} e_i = 1\right) \cdot \Pr(e_m = 0) + \Pr\left(\sum_{i=1}^{m-1} e_i = 0\right) \cdot \Pr(e_m = 1) \\ &= (1 - \tau)\left(\frac{1}{2} - \frac{1}{2}(1 - 2\tau)^{m-1}\right) + \tau\left(\frac{1}{2} + \frac{1}{2}(1 - 2\tau)^{m-1}\right) \\ &= \frac{1}{2} - \frac{1}{2}(1 - 2\tau)^m \end{aligned}$$

□

Remark. In the naive Gaussian elimination, we sum up $\mathcal{O}(n^2)$ LPN $_n^\tau$ samples, with bias $\eta = \frac{1}{2} - \tau$, which results in LPN samples with bias $\eta' = \frac{1}{2}(1 - 2\tau)^{n^2} = c^{-n^2}$. One would need $m \geq c^{n^2}$ many such samples to decide on S_1

3.3 The [BKW00] algorithm

The algorithm consists of two parts:

1. Block Gaussian elimination (aim: produce $e_1 = \sum_{i \in I} a_i$, $|I| = n^{1-\epsilon}$)
2. Majority vote

Algorithm 4: BKW part 1

Input: a list $\mathcal{L}^{(0)} = \mathcal{L} = \{(a_i, b_i) \text{ LPN samples}\}$ of size m
 $k \in \mathbb{Z}$ block size (assume $k|n$)

Output: $I \subseteq [m]$ s.t. $|I| = 2^{\frac{n}{k}}$ and $\sum_{i \in I} a_i = e_1$

- 1: **for** $i \in [0, \frac{n}{k} - 2]$ **do**
 - 2: $\mathcal{L}^{(i+1)} \leftarrow \{\}$
 - 3: Sort $\mathcal{L}^{(i)}$ w.r.t last k nonzero coordinates
 - 4: **for all** $j \in \{0, 1\}^k$ **do**
 - 5: choose an element $a_j \in \mathcal{L}^{(i)}$ whose last k nonzero coordinates are equal to j
 - 6: $\mathcal{L}^{(i+1)} = \mathcal{L}^{(i+1)} \cup \{(a_j + a'_j, b_j + b'_j) \text{ for all } a'_j \neq a_j \text{ whose last nonzero coordinates are } = j\}$
 - 7: **end for**
 - 8: **end for**
 - 9: Sort $\mathcal{L}^{(\frac{n}{k}-2)}$ w.r.t the last $k - 1$ coordinates.
 - 10: Find a pair of elements (x_1, x_2) from $\mathcal{L}^{(\frac{n}{k}-2)}$ s.t. $x_1[1] + x_2[1] = e_1$
-

Claim 8. Taking $m = |\mathcal{L}| = \Theta(\text{poly}(n) \cdot 2^k)$ suffices for Algorithm 4 to output $\mathcal{L}^{(\frac{n}{k}-1)}$ of expected size 1 in time $T(\text{Algorithm 4}) = \tilde{\mathcal{O}}(2^k)$ using memory $M(\text{Algorithm 4}) = \tilde{\mathcal{O}}(2^k)$.

Proof sketch. On each step i , we partition $\mathcal{L}^{(i)}$ into at most 2^k classes represented by j . For each non-empty class (i.e. \exists at least two a_j, a'_j in $\mathcal{L}^{(i)}$), we discard only one of its representatives. $\implies |\mathcal{L}^{(i+1)}| > |\mathcal{L}^{(i)}| - 2^k$ provided all (almost all) classes are non-empty. We need $\Omega(\text{poly}(k) \cdot 2^k)$ elements in $\mathcal{L}^{(i)}$ for $(1 - e^{-n})$ -fraction of all classes to contain at least two elements. \implies we need to start with $|\mathcal{L}^{(0)}| = \Omega(\frac{n}{k} \text{poly}(k) \cdot 2^k) = \Omega(\text{poly}(n) \cdot 2^k)$ elements. Elements in $\mathcal{L}^{(i)}$ are uniformly random, conditioned on having zeros on the last $(i - k)$ coordinates. Indeed, let us denote $Y = |\mathcal{L}^{(i)}|$ the random variable giving the size of $\mathcal{L}^{(i)}$. Cf Markov, $\Pr(Y > a) \leq \frac{\mathbb{E}(Y)}{a}$. Taking $a = \mathbb{E}(Y) \cdot \text{poly}(n)$ gives us a $1 - \frac{1}{\text{poly}(n)}$ probability on the list bound, and hence a $\frac{1}{\text{poly}(n)}$ probability on the success of the algorithm. (We could take $a = \mathbb{E}(Y) \cdot 2^{\epsilon n}$ to have an overwhelming probability. In that case, we need to replace $\tilde{\mathcal{O}}(2^k)$ in the claim by $2^{\mathcal{O}(k)}$. This will however not change the next theorem). See [DRX17] for more details.

The most expensive part of the algorithm is sorting, it takes time $\tilde{\mathcal{O}}(|\mathcal{L}^{(i)}|)$. \square

Algorithm 5: The BKW algorithm

Input: m -many LPN_n^τ samples, $\epsilon > 0$

Output: $s \in \{0, 1\}^n$

- 1: **for** $i \in [n]$ **do**
 - 2: Run Algorithm 4 with $k = \frac{n}{(1-\epsilon)\log(n)}$, $N = \Theta((1 - 2\tau)^{n^{1-\epsilon}})$ times to obtain $(e_i, b_j)_{j \leq N}$.
 - 3: Run the majority vote algorithm to decide on s_i .
 - 4: **end for**
-

Theorem 9. The BKW algorithm 5 solves the LPN_n^τ problem in time $T(\text{BKW}) = 2^{\mathcal{O}(\frac{n}{\log(n)})}$ using $M(\text{BKW}) = 2^{\mathcal{O}(\frac{n}{\log(n)})}$ memory and LPN samples.

Proof. Algorithm 4 will output an $LPN_n^{\tau'}$ sample of the form (e_i, b_i) in time $\tilde{O}(2^{\frac{n}{(1-\epsilon)\log(n)}})$, with $\tau' = \frac{1}{2} - \frac{1}{2}(1-2\tau)^{2^{(1-\epsilon)\log(n)}} = \frac{1}{2} - \frac{1}{2}C^{n^{1-\epsilon}}$ with C a constant depending on τ .

For the majority vote, we would need to repeat Algorithm 4 $\mathcal{O}(2^{n^{1-\epsilon}})$ times to decide on s_i with constant success probability. \square

Remark. 1. A more precise complexity of BKW is:

$$T = M = \#samples = 2^{\frac{n}{\log(\frac{n}{\tau})}(1+o(1))} \text{ cf [EKM17]}$$

2. Given $m = n^{1+\epsilon}$ many LPN_n^τ samples, the BKW algorithm's run-time goes up to $T(BKW) = 2^{bigO(\frac{n}{\log\log(n)})}$
(sample "amplification", cf V. Lyubashevsky in [Lyu05])
3. The "LF2" technique due to Levieil-Fouque [LF06] analyses the BKW algorithm by considering all possible pairs during the zeroizing step.
Proved by Devadas et al. [DRX17].

Ring-LPN

Take the ring $\mathbb{F}_2[X]/(f)$, with f a degree n polynomial.

Ring-LPN sample $s \in \mathbb{F}_2[X]/(f)$. ($a_i \xleftarrow{\$} \mathbb{F}_2[X]/(f)$, $a_i \cdot s_i + e_i \in \mathbb{F}_2[X]/(f)$).
 e_i : polynomial with coefficients chosen from Ber_τ .

Open problem: find a better algorithm for Ring-LPN than as for "standard" LPN.

References

- [Ale03] Michael Alekhnovich. More on average case vs approximation complexity. In *Proceedings of the 44th Annual IEEE Symposium on Foundations of Computer Science*, FOCS '03, pages 298–, Washington, DC, USA, 2003. IEEE Computer Society.
- [BKW00] Avrim Blum, Adam Kalai, and Hal Wasserman. Noise-tolerant learning, the parity problem, and the statistical query model. *CoRR*, cs.LG/0010022, 2000.
- [DRX17] Srinivas Devadas, Ling Ren, and Hanshen Xiao. On iterative collision search for lpn and subset sum. Cryptology ePrint Archive, Report 2017/904, 2017. <https://eprint.iacr.org/2017/904>.
- [EKM17] Andre Esser, Robert Kbler, and Alexander May. Lpn decoded. Cryptology ePrint Archive, Report 2017/078, 2017. <https://eprint.iacr.org/2017/078>.
- [HB01] Nicholas J. Hopper and Manuel Blum. Secure human identification protocols. In Colin Boyd, editor, *Advances in Cryptology — ASIACRYPT 2001*, pages 52–66, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- [LF06] Éric Levieil and Pierre-Alain Fouque. An improved lpn algorithm. In Roberto De Prisco and Moti Yung, editors, *Security and Cryptography for Networks*, pages 348–359, Berlin, Heidelberg, 2006. Springer Berlin Heidelberg.

- [Lyu05] Vadim Lyubashevsky. The parity problem in the presence of noise, decoding random linear codes, and the subset sum problem. In *Proceedings of the 8th International Workshop on Approximation, Randomization and Combinatorial Optimization Problems, and Proceedings of the 9th International Conference on Randomization and Computation: Algorithms and Techniques*, APPROX'05/RANDOM'05, pages 378–389, Berlin, Heidelberg, 2005. Springer-Verlag.
- [MU05] Michael Mitzenmacher and Eli Upfal. *Probability and Computing: Randomized Algorithms and Probabilistic Analysis*. Cambridge University Press, Cambridge, 2005.